

Project_Team1

May 21, 2018

Introduction to Machine Learning (Spring 18) - Team 1
Team Members: Anurag Marwah (Net ID: am8482)

1 Initialize

```
In [1]: import numpy as np
        from scipy import misc
        import time
        import matplotlib.pyplot as plt
        import pandas as pd
        %matplotlib inline
```

2 Step 0: Prepare Data

2.1 Flickr API

```
In [2]: import flickrapi
        import urllib.request
        import skimage.io
        import skimage.transform
        import requests
        import warnings
        import os
        from io import BytesIO

In [3]: api_key = u'xxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxx'
        api_secret = u'xxxxxxxxxxxxxxxx'
        flickr = flickrapi.FlickrAPI(api_key, api_secret)

In [4]: def createdir(dir_name):
        dir_exists = os.path.isdir(dir_name)
        if not dir_exists:
            os.mkdir(dir_name)
            print("Making directory %s" % dir_name)
        else:
            print("Directory %s already present!" % dir_name)
```

```

In [5]: def downloadimages(dir_name,keyword,nimage):

    photos = flickr.walk(text=keyword, tag_mode='all', tags=keyword,extras='url_c',\
                          sort='relevance',per_page=100)

    #     nimage = 10
    i = 0
    nrow = 224
    ncol = 224
    for photo in photos:
        url=photo.get('url_c')
        if not (url is None):

            # Create a file from the URL
            # This may only work in Python3
            response = requests.get(url)
            file = BytesIO(response.content)

            # Read image from file
            im = skimage.io.imread(file)

            # Resize images
            im1 = skimage.transform.resize(im,(nrow,ncol),mode='constant')

            # Convert to uint8, suppress the warning about the precision loss
            with warnings.catch_warnings():
                warnings.simplefilter("ignore")
                im2 = skimage.img_as_ubyte(im1)

            # Save the image
            local_name = '{0:s}/{1:s}/{2:s}_{3:04d}.jpg'.format(dir_name,keyword,keyword,i)
            skimage.io.imsave(local_name, im2)
            print(local_name)
            i = i + 1
    if (i >= nimage):
        break

```

2.2 Creating Directories

```

In [6]: dir1 = 'train'
        dir2 = 'test'
        keyword1 = 'person'
        keyword2 = 'car'
        keyword3 = 'animal'
        keyword4 = 'bird'
        keyword5 = 'building'
        keyword6 = 'tree'
        keyword7 = 'face'
        keyword8 = 'group'

```

```
In [7]: # createdir(dir1)
# createdir(dir1+'/'+keyword1)
# createdir(dir1+'/'+keyword2)
# createdir(dir1+'/'+keyword3)
# createdir(dir1+'/'+keyword4)
# createdir(dir1+'/'+keyword5)
# createdir(dir1+'/'+keyword6)
# createdir(dir1+'/'+keyword7)
# createdir(dir2)
# createdir(dir2+'/'+keyword1)
# createdir(dir2+'/'+keyword2)
# createdir(dir2+'/'+keyword3)
# createdir(dir2+'/'+keyword4)
# createdir(dir2+'/'+keyword5)
# createdir(dir2+'/'+keyword6)
# createdir(dir2+'/'+keyword7)
# createdir(dir2+'/'+keyword8)
```

2.3 Download Data

```
In [8]: # downloadimages(dir1,keyword1,1000)
# downloadimages(dir1,keyword2,1000)
# downloadimages(dir1,keyword3,1000)
# downloadimages(dir1,keyword4,1000)
# downloadimages(dir1,keyword5,1000)
# downloadimages(dir1,keyword6,1000)
# downloadimages(dir1,keyword7,1000)

# downloadimages(dir2,keyword1,300)
# downloadimages(dir2,keyword2,300)
# downloadimages(dir2,keyword3,300)
# downloadimages(dir2,keyword4,300)
# downloadimages(dir2,keyword5,300)
# downloadimages(dir2,keyword6,300)
# downloadimages(dir2,keyword7,300)
# downloadimages(dir2,keyword8,300)
```

3 Step 1: Download Model

```
In [9]: import keras
```

/Applications/anaconda/lib/python3.6/site-packages/h5py/__init__.py:34: FutureWarning: Convers
 from ._conv import register_converters as _register_converters
 Using TensorFlow backend.

```
In [10]: from keras import applications
from keras.preprocessing.image import ImageDataGenerator
```

```

from keras import optimizers
from keras.models import Sequential
from keras.layers import Dropout, Flatten, Dense

```

```

In [11]: import keras.backend as K
        K.clear_session()

```

```

In [12]: nrow = 64
        ncol = 64
        nchan = 3
        image_shape = (nrow,ncol,3)
        print(image_shape)

```

(64, 64, 3)

3.1 Download VGGNet

```

In [13]: from keras.applications.vgg16 import VGG16
        from keras.preprocessing import image
        from keras.applications.vgg16 import decode_predictions, preprocess_input
        base_model = VGG16(include_top=False, weights='imagenet', input_shape=image_shape)
        base_model.summary()

```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 64, 64, 3)	0
block1_conv1 (Conv2D)	(None, 64, 64, 64)	1792
block1_conv2 (Conv2D)	(None, 64, 64, 64)	36928
block1_pool (MaxPooling2D)	(None, 32, 32, 64)	0
block2_conv1 (Conv2D)	(None, 32, 32, 128)	73856
block2_conv2 (Conv2D)	(None, 32, 32, 128)	147584
block2_pool (MaxPooling2D)	(None, 16, 16, 128)	0
block3_conv1 (Conv2D)	(None, 16, 16, 256)	295168
block3_conv2 (Conv2D)	(None, 16, 16, 256)	590080
block3_conv3 (Conv2D)	(None, 16, 16, 256)	590080
block3_pool (MaxPooling2D)	(None, 8, 8, 256)	0

block4_conv1 (Conv2D)	(None, 8, 8, 512)	1180160

block4_conv2 (Conv2D)	(None, 8, 8, 512)	2359808

block4_conv3 (Conv2D)	(None, 8, 8, 512)	2359808

block4_pool (MaxPooling2D)	(None, 4, 4, 512)	0

block5_conv1 (Conv2D)	(None, 4, 4, 512)	2359808

block5_conv2 (Conv2D)	(None, 4, 4, 512)	2359808

block5_conv3 (Conv2D)	(None, 4, 4, 512)	2359808

block5_pool (MaxPooling2D)	(None, 2, 2, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

3.2 Create New Model

```
In [14]: # Create new model
        model = Sequential()

        # Add layers into the new model
        for layer in base_model.layers:
            model.add(layer)

        # Freeze the layers in the new model
        for layer in model.layers:
            layer.trainable=False

        model.summary()
```

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	(None, 64, 64, 3)	0

block1_conv1 (Conv2D)	(None, 64, 64, 64)	1792

block1_conv2 (Conv2D)	(None, 64, 64, 64)	36928

block1_pool (MaxPooling2D)	(None, 32, 32, 64)	0

block2_conv1 (Conv2D)	(None, 32, 32, 128)	73856
block2_conv2 (Conv2D)	(None, 32, 32, 128)	147584
block2_pool (MaxPooling2D)	(None, 16, 16, 128)	0
block3_conv1 (Conv2D)	(None, 16, 16, 256)	295168
block3_conv2 (Conv2D)	(None, 16, 16, 256)	590080
block3_conv3 (Conv2D)	(None, 16, 16, 256)	590080
block3_pool (MaxPooling2D)	(None, 8, 8, 256)	0
block4_conv1 (Conv2D)	(None, 8, 8, 512)	1180160
block4_conv2 (Conv2D)	(None, 8, 8, 512)	2359808
block4_conv3 (Conv2D)	(None, 8, 8, 512)	2359808
block4_pool (MaxPooling2D)	(None, 4, 4, 512)	0
block5_conv1 (Conv2D)	(None, 4, 4, 512)	2359808
block5_conv2 (Conv2D)	(None, 4, 4, 512)	2359808
block5_conv3 (Conv2D)	(None, 4, 4, 512)	2359808
block5_pool (MaxPooling2D)	(None, 2, 2, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 0		
Non-trainable params: 14,714,688		

```
In [15]: model.add(Flatten())
         model.add(Dense(256, activation='relu'))
         model.add(Dense(1, activation='sigmoid'))
         model.summary()
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 64, 64, 3)	0
block1_conv1 (Conv2D)	(None, 64, 64, 64)	1792

block1_conv2 (Conv2D)	(None, 64, 64, 64)	36928

block1_pool (MaxPooling2D)	(None, 32, 32, 64)	0

block2_conv1 (Conv2D)	(None, 32, 32, 128)	73856

block2_conv2 (Conv2D)	(None, 32, 32, 128)	147584

block2_pool (MaxPooling2D)	(None, 16, 16, 128)	0

block3_conv1 (Conv2D)	(None, 16, 16, 256)	295168

block3_conv2 (Conv2D)	(None, 16, 16, 256)	590080

block3_conv3 (Conv2D)	(None, 16, 16, 256)	590080

block3_pool (MaxPooling2D)	(None, 8, 8, 256)	0

block4_conv1 (Conv2D)	(None, 8, 8, 512)	1180160

block4_conv2 (Conv2D)	(None, 8, 8, 512)	2359808

block4_conv3 (Conv2D)	(None, 8, 8, 512)	2359808

block4_pool (MaxPooling2D)	(None, 4, 4, 512)	0

block5_conv1 (Conv2D)	(None, 4, 4, 512)	2359808

block5_conv2 (Conv2D)	(None, 4, 4, 512)	2359808

block5_conv3 (Conv2D)	(None, 4, 4, 512)	2359808

block5_pool (MaxPooling2D)	(None, 2, 2, 512)	0

flatten_1 (Flatten)	(None, 2048)	0

dense_1 (Dense)	(None, 256)	524544

dense_2 (Dense)	(None, 1)	257
=====		
Total params: 15,239,489		
Trainable params: 524,801		
Non-trainable params: 14,714,688		

4 Step 2: Train Model

4.1 Data Generators

```
In [16]: train_data_dir = './train'
        batch_size = 32
        train_datagen = ImageDataGenerator(rescale=1./255,
                                           shear_range=0.2,
                                           zoom_range=0.2,
                                           horizontal_flip=True)
        train_generator = train_datagen.flow_from_directory(
            train_data_dir,
            target_size=(nrow,ncol),
            batch_size=batch_size,
            classes=['notaface', 'face'],
            class_mode='binary')
```

Found 4000 images belonging to 2 classes.

```
In [17]: test_data_dir = './test'
        batch_size = 32
        test_datagen = ImageDataGenerator(rescale=1./255,
                                           shear_range=0.2,
                                           zoom_range=0.2,
                                           horizontal_flip=True)
        test_generator = test_datagen.flow_from_directory(
            test_data_dir,
            target_size=(nrow,ncol),
            batch_size=batch_size,
            classes=['notaface', 'face'],
            class_mode='binary')
```

Found 2000 images belonging to 2 classes.

```
In [18]: # Display the image
        def disp_image(im):
            if (len(im.shape) == 2):
                # Gray scale image
                plt.imshow(im, cmap='gray')
            else:
                # Color image.
                im1 = (im-np.min(im))/(np.max(im)-np.min(im))*255
                im1 = im1.astype(np.uint8)
                plt.imshow(im1)


        # Remove axis ticks
        plt.xticks([])
        plt.yticks([])
```


4.2 Sample Train and Test Data

```
In [19]: nplot = 10
plt.figure(figsize=(20,20))
minibatch_x,minibatch_y = train_generator.next()
for i in range(nplot):
    plt.subplot(1,nplot,i+1)
    disp_image(minibatch_x[i])
    plt.title(minibatch_y[i])
```



```
In [20]: nplot = 10
plt.figure(figsize=(20,20))
minibatch_x,minibatch_y = test_generator.next()
for i in range(nplot):
    plt.subplot(1,nplot,i+1)
    disp_image(minibatch_x[i])
    plt.title(minibatch_y[i])
```



4.3 Model Training

```
In [21]: model.compile(loss='binary_crossentropy',
                        optimizer=optimizers.adam(lr=1e-3),
                        metrics=['accuracy'])
```

```
In [22]: steps_per_epoch = train_generator.n // batch_size
validation_steps = test_generator.n // batch_size
```

```
In [23]: nepochs = 5 # Number of epochs

# Call the fit_generator function
hist = model.fit_generator(
    train_generator,
    steps_per_epoch=steps_per_epoch,
    epochs=nepochs,
    validation_data=test_generator,
    validation_steps=validation_steps)
```

```

Epoch 1/5
125/125 [=====] - 320s 3s/step - loss: 0.0447 - acc: 0.9818 - val_loss: 0.0447
Epoch 2/5
125/125 [=====] - 327s 3s/step - loss: 0.0073 - acc: 0.9982 - val_loss: 0.0073
Epoch 3/5
125/125 [=====] - 316s 3s/step - loss: 0.0117 - acc: 0.9955 - val_loss: 0.0117
Epoch 4/5
125/125 [=====] - 324s 3s/step - loss: 0.0033 - acc: 0.9985 - val_loss: 0.0033
Epoch 5/5
125/125 [=====] - 327s 3s/step - loss: 0.0046 - acc: 0.9982 - val_loss: 0.0046

```

4.4 Training Accuracy vs Validation Accuracy - Plot

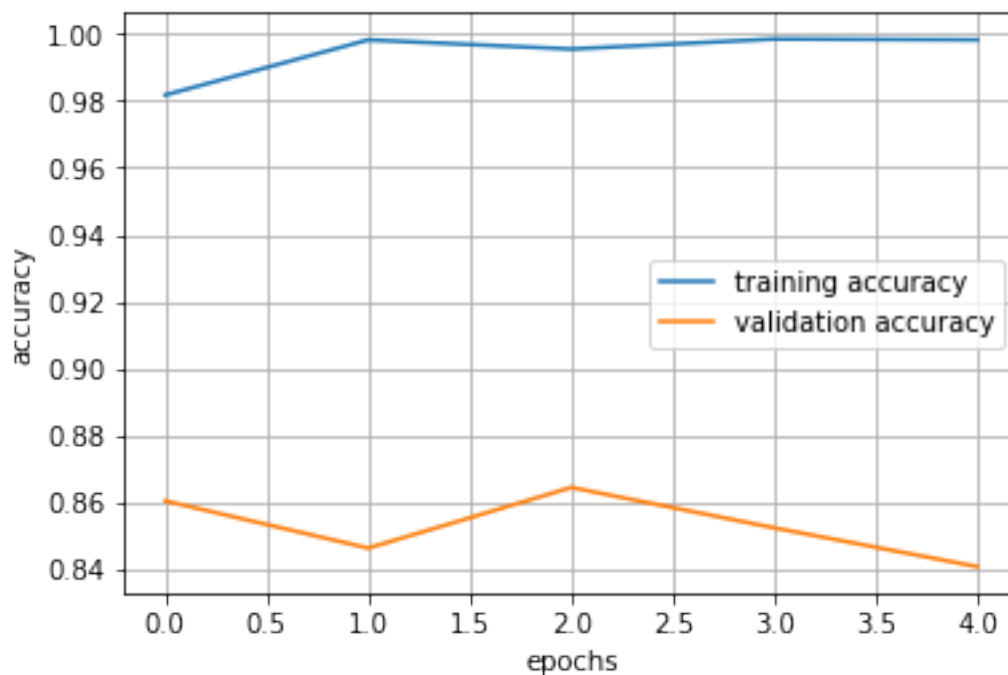
In [24]: *# Plot the training accuracy and validation accuracy curves on the same figure.*

```

tr_accuracy = hist.history['acc']
val_accuracy = hist.history['val_acc']
plt.plot(tr_accuracy)
plt.plot(val_accuracy)
plt.grid()
plt.xlabel('epochs')
plt.ylabel('accuracy')
plt.legend(['training accuracy', 'validation accuracy'])

```

Out[24]: <matplotlib.legend.Legend at 0x122229f28>



5 Step 3: Predict Faces

```
In [25]: from skimage.transform import pyramid_gaussian
```

```
In [26]: import matplotlib.patches as patches
```

5.1 Sliding Window

```
In [27]: def sliding_window(image, stepSize, windowSize):
    # slide a window across the image
    for y in range(0, image.shape[0], stepSize):
        for x in range(0, image.shape[1], stepSize):
            # yield the current window
            yield (x, y, image[y:y + windowSize[1], x:x + windowSize[0]])

In [28]: # Display the image
def disp_image2(ax,im):
    if (len(im.shape) == 2):
        # Gray scale image
        ax.imshow(im, cmap='gray')
    else:
        im1 = (im-np.min(im))/(np.max(im)-np.min(im))*255
        im1 = im1.astype(np.uint8)
        ax.imshow(im1)

    # Remove axis ticks
    plt.xticks([])
    plt.yticks([])
```

5.2 Face Prediction

```
In [32]: def predict_face(model,im,scale_factor,winW,winH,step,threshold):

    # Build Image Pyramid
    pyramid = tuple(pyramid_gaussian(im, downscale=scale_factor))

    # Calculate the number of images to be used from the image pyramid
    for pyramid_count,p in enumerate(pyramid):
        # if the image is too small, break from the loop
        if p.shape[0] < 64 or p.shape[1] < 64:
            break

    # Build Image Plot
    fig,ax = plt.subplots(1)
    disp_image2(ax,im)

    # Iterate over the Image Pyramid
    for j in range(pyramid_count):
        p = pyramid[j]
```

```

# Get windows
i=0
x_win = ((int)((p.shape[1]-winW)/step))+1
y_win = ((int)((p.shape[0]-winH)/step))+1
num_boxes = x_win * y_win
box = np.zeros((num_boxes,2))
win_arr = np.zeros((num_boxes,winW,winH,3))

# Iterate over the Sliding Windows
for (x, y, window) in sliding_window(p, stepSize=step, \
                                     windowSize=(winW, winH)):

    # if the window does not meet our desired window size, ignore it
    if window.shape[0] != winH or window.shape[1] != winW:
        continue

    # Get windows in an array
    win_arr[i] = window
    box[i] = x,y
    i = i + 1

# Use VGG-Net for Image Class Prediction
ntest = win_arr.shape[0]
yhat = model.predict(win_arr)
yhat_bin = (yhat>threshold).reshape(ntest).astype(int)

# Print Boxes
for k,c in enumerate(box):
    if yhat_bin[k]==1:
        rect = patches.Rectangle((box[k,0],box[k,1]),\
                                winW*(np.power(scale_factor,j)),\
                                winH*(np.power(scale_factor,j)),\
                                fill=False,ls='solid',lw=1,color='b')

        ax.add_patch(rect)

```


```

In [30]: im_w = 360
         im_h = 270
         im_shape = (im_h,im_w,3)
         im = np.zeros(im_shape)
         class_name = 'group'
         fn = '{0:s}/{1:s}_{2:04d}.jpg'.format(class_name,class_name,3)
         im[:, :, :] = image.load_img(fn, target_size=(im_h, im_w))
         disp_image(im)

```



```
In [33]: predict_face(model,im,scale_factor = 2,winW = 64,winH = 64,step = 8,threshold = 0.5)
```



```
In [ ]:
```